

# Introduction to Music Information Retrieval using Essentia.js

Albin Correy, Dmitry Bogdanov, Luis Joglar-Ongay\*, Jorge Marcos-Fernández

Music Technology Group, Universitat Pompeu Fabra

\*SonoSuite

<https://github.com/MTG/essentia.js-tutorial-wac2021>

Web Audio Conference 2021

# About us



**Albin Correyea**

[@albincorreyea](#)

Research collaborator at **MTG**  
Senior ML Engineer at  
**Moodagent A/S, Denmark**  
Dev of Essentia.js



**Dmitry Bogdanov**

[@di\\_bogdanov](#)  
<https://dbogdanov.com>

Senior researcher and  
lead developer of Essentia  
at **MTG UPF**



**Jorge Marcos**

[@MaferGeorge](#)

Researcher and developer at  
**MTG UPF**  
Working on Essentia.js  
applications



**Luis Joglar-Ongay**

[@luisjoglar](#)

PhD Candidate at **MTG UPF**  
Research engineer at **SonoSuite**  
Collaborator of Essentia.js



Universitat  
Pompeu Fabra  
Barcelona

**MTG**  
Music Technology  
Group

<https://www.upf.edu/web/mtg/>

# About this tutorial

1. Introduction to MIR and audio analysis. MIR applications and a typical analysis pipeline.
2. Using Essentia.js for music and audio analysis. Overview of available music audio features and application use-cases.
3. Getting started with Essentia.js. Writing your first “Hello world” application. Using Essentia.js for deferred-time vs. real-time analysis.
4. Available demos and template projects in a JavaScript playground.
5. Deep learning inference with Essentia.js using pre-trained machine learning models. Interface with TensorFlow.js.
6. Demos and examples of using Essentia.js for machine learning inference.
7. An industrial use-case example: audio problem detection for music distribution with Essentia.js.
8. QA and a playground session.

# Audio analysis and Music Information Retrieval (MIR)

- Audio analysis (Wikipedia): *“Extraction of information and meaning from audio signals for analysis, classification, storage, retrieval, and synthesis.”*
- Since 2000, we have a dedicated conference:  
[International Society for Music Information Retrieval \(ISMIR\) conference](#)
- ISMIR web page: *“processing, searching, organising and accessing music-related data.”*
- Wikipedia: *“the interdisciplinary science of retrieving information from music”*.
- Lots of MIR research is audio-based
- Good ISMIR tutorials:
  - [A Basic Introduction to Audio-Related Music Information Retrieval](#)
  - <https://www.upf.edu/web/mtg/mir-overview-recent-developments-future>
- Similarly, there is a lot of research on **sound information retrieval**

# Some examples of music and sound retrieval tasks

Music/sound browsing interfaces

Music/sound search, exploration and recommendation

Music/sound classification

Semantic/tag-based retrieval

Query by example

Query by humming

Music source separation and instrument recognition

Music identification, cover song identification

Automatic music transcription

Music alignment (e.g., audio-to-score, audio-to-lyrics)

Music/sound generation, creative applications

Music/sound visualization

# Extracting audio features

Many possibilities for applications operating with music/sound features retrieved from audio. Extracting this features is a fundamental step.

## Low- and mid-level **sound and music features**

sound envelope, loudness, timbre, tonality, pitch/melody, onsets, rhythm, etc.

## High-level **semantic descriptors**

type of sound, music genres, instrumentation, moods, danceability, etc.

# Essentia <https://essentia.upf.edu>

Open-source library and tools for audio and music analysis, description and synthesis

- Extensive collection of reusable algorithms
- Written in **C++** and optimized for computational speed
- **Python** bindings for fast prototyping
- Feature extractors for **large-scale audio analysis**
- **Cross-platform** (Linux, Mac OS X, Windows, iOS, Android)
- Support for mobile platforms and **real-time** processing
- Developed at Music Technology Group, UPF

```
1 from essentia.standard import *
2 audio = MonoLoader(filename='audio.mp3')()
3 beats, bconfidence = BeatTrackerMultiFeature()(audio)
4 audio = EqualLoudness()(audio)
5 melody, mconfidence = PitchMelodia(frameSize=2048, hopSize=128)(audio)
```

# Essentia

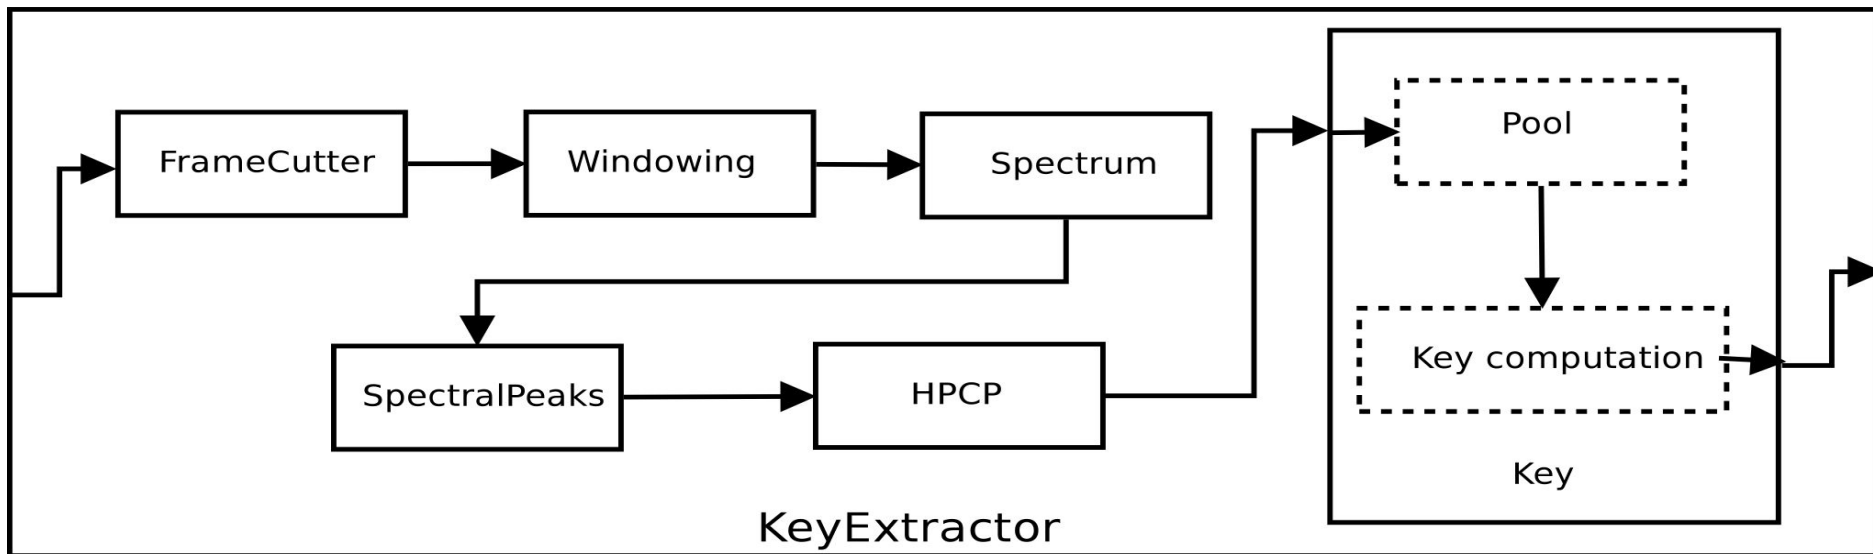
260+ algorithms for audio signal processing and analysis, and sound and music description, developed at MTG

- Standard audio IO & DSP
- Sound and music descriptors
  - spectral features
  - rhythm and tempo
  - tonality, pitch and melody
  - loudness/dynamics
  - sound envelope
  - audio segmentation
  - fingerprinting
- Machine-learning based descriptors
  - genres, moods, instrumentation, ...
  - SVM classifiers, inference with TensorFlow deep learning models



# Modularity and processing chains

Users can build their own extractors for the descriptors they want to compute in a “data-flow” manner



# Many applications in MIR and beyond





# Music and Audio analysis on the Web, so far

Plenty of web applications use audio analysis capabilities

But ...

Most of them do audio analysis on the servers

Since

- Limitations of the client devices and the web browsers itself.
- Web Audio API capabilities are limited for audio analysis.
- Lack of extensive software libraries for on-client computation.

# Existing audio analysis APIs and libraries on the web

Spotify API - audio features for music in Spotify

Freesound API - audio analysis for sounds in Freesound with Essentia

AcousticBrainz API - database of audio features extracted with Essentia

Meyda <https://meyda.js.org> - written purely in JS

JS-Xtract <https://github.com/nickjillings/js-xtract> - JS conversion of LibXtract

aubiojs <https://qiuxiang.github.io/aubiojs> - JS conversion of some Aubio algorithms

# Music/audio analysis with Essentia and Essentia.js

The logo for Essentia, featuring a red square icon with a white stylized 'E' and the word 'SENTIA' in a dark blue, sans-serif font.


C++ library

<https://essentia.upf.edu>



An open-source JavaScript library,  
powered by Essentia WebAssembly

<https://essentia.upf.edu/essentiajs>



Extensive collection of  
reusable algorithms

Flexible and easily extendable  
algorithms for common audio  
analysis processes and audio and  
music descriptors.

Cross-platform

Linux, Mac OS X, Windows, iOS,  
Android, and Web.

Fast prototyping

Python scientific environment,  
JavaScript bindings, and  
command-line audio analysis  
tools.

Industrial applications

Optimized for computational  
speed, including real-time use  
cases.



An open-source Javascript library for music/audio analysis and  
processing,  
powered by WebAssembly

<https://essentia.upf.edu/essentiajs>

Correya, A. A., Bogdanov, D., Joglar-Ongay, L., & Serra, X. (2020). Essentia. js: a JavaScript library for music and audio analysis on the web. *Proceedings of the 21st International Society for Music Information Retrieval Conference; 2020 Oct 11-16; Montréal, Canada.[Canada]; ISMIR; 2020. p. 605-12.. International Society for Music Information Retrieval (ISMIR).*

# License

**AGPLv3** + commercial license under request



<https://www.gnu.org/licenses/agpl-3.0.en.html>

<https://opensource.stackexchange.com/questions/4442/how-does-the-agpl-apply-to-javascript-libraries>



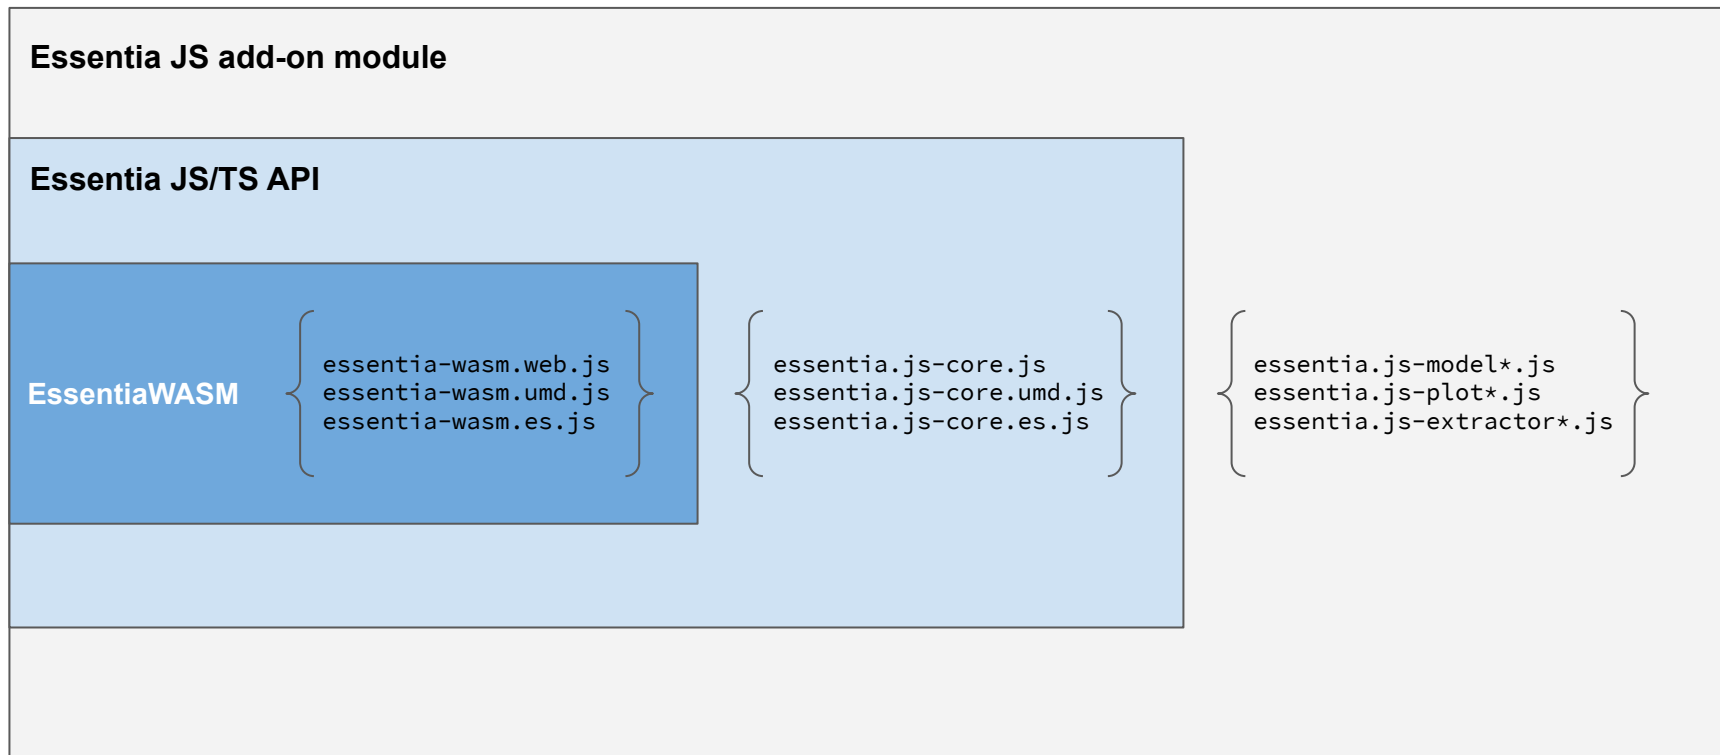
# Let's checkout `essentia.js` in action!

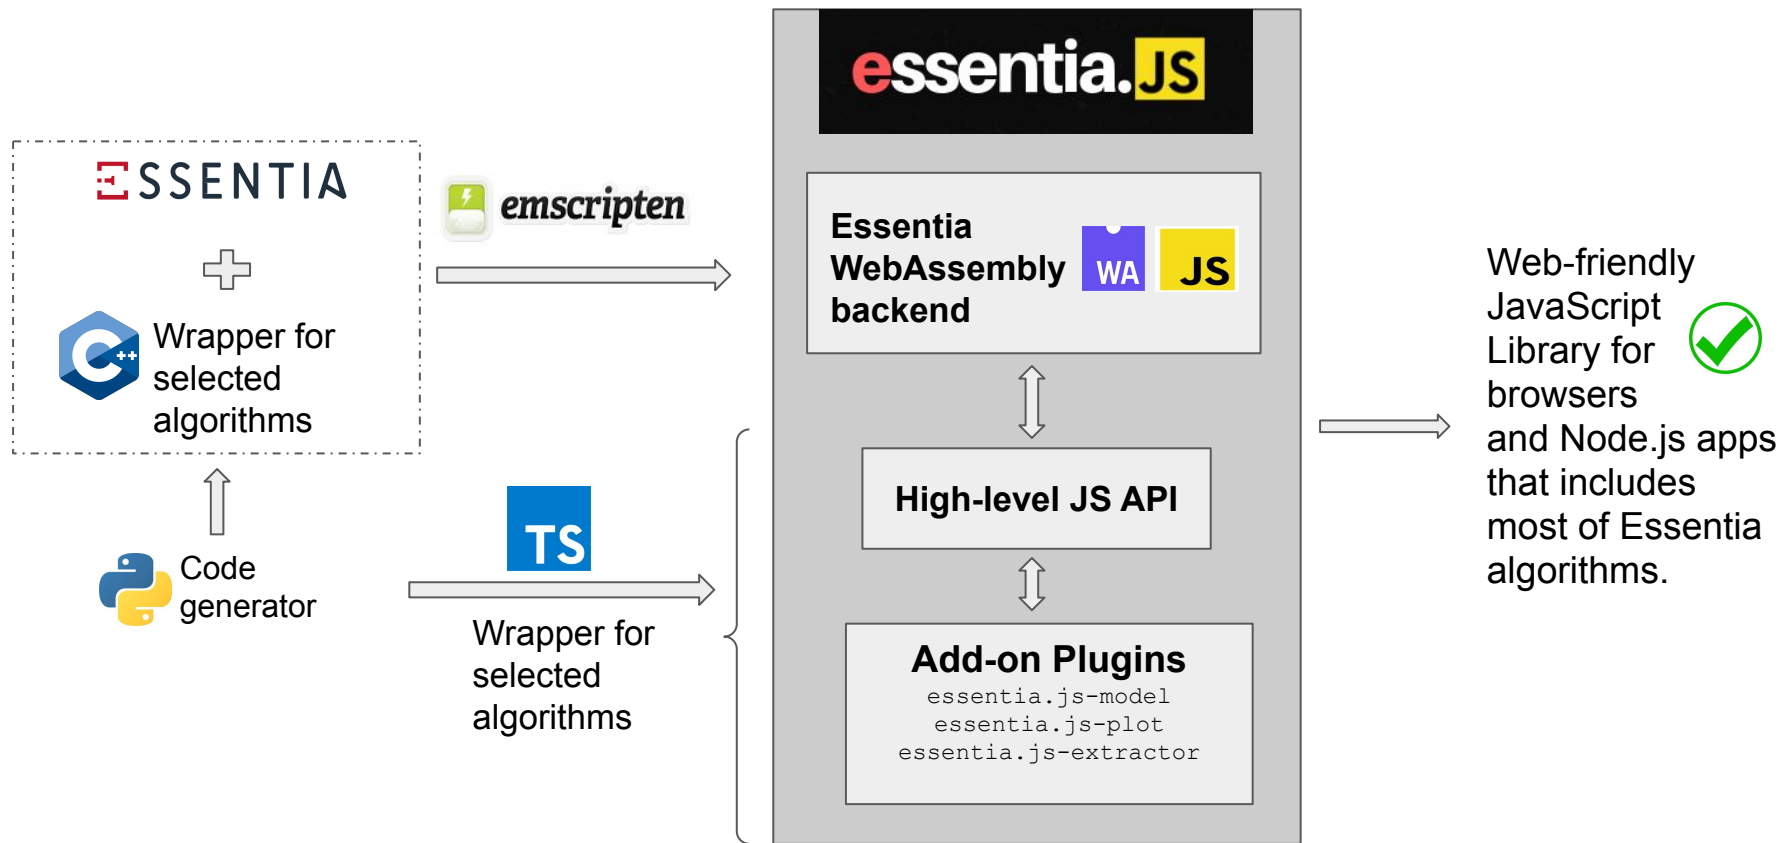
[essentia.js examples](#)

# Design choices & functionalities

- Easy installation
  - NPM, CDN etc
- User friendly API & utility tools
  - Easy-to-use
- Modularity & Extensibility
  - Easy-to-customize and build
- Web standards compliance
  - Web Audio API, Audio Worklet, ES6 etc.
- Lightweight and lesser dependencies
  - Small as 2.7 MB ~
- Reproducibility across different platforms
  - Native and Web (C++, Python and JavaScript)
- Extensive documentation and examples

# Abstractions





# Known limitations

- As of now, only 200 essentia algorithms are supported ([here](#)).
- But all algorithms are available via [custom C++ wrapper](#).
- Essentia.js has not been through rigorous QA tests. Hence, not fully production ready and backward compatible.

But, we are currently on active development on making towards a stable production-ready release.

You're most welcome to contribute!

# Getting started



```
npm install essentia.js
```



**script** `</>`

```
<script src="https://cdn.jsdelivr.net/npm/essentia.js@0.1.1/dist/essentia-wasm.web.js"></script>  
<script src="https://cdn.jsdelivr.net/npm/essentia.js@0.1.1/dist/essentia.js-model.js"></script>
```

## ES6 style imports

```
import { EssentiaWASM } from 'https://cdn.jsdelivr.net/npm/essentia.js@0.1.1/dist/essentia-wasm.module.js';  
import * as EssentiaModel from 'https://cdn.jsdelivr.net/npm/essentia.js@0.1.1/dist/essentia.js-model.js';
```

# Basic HTML use

Steps:

Import the library

Instantiate Essentia.js with WASM backend

Get audio

Fetch file and decode

Request live microphone stream \*

Use Essentia.js algorithms

```
<script
  src="https://cdn.jsdelivr.net/npm/essentia.js@0.1.3/dist/essentia-wasm.web.js"
  defer>
</script>
<script
  src="https://cdn.jsdelivr.net/npm/essentia.js@0.1.3/dist/essentia.js-core.js"
  defer>
</script>
```

```
let essentia = null;
window.onload = () => {
  // load Essentia WASM backend
  EssentiaWASM().then(wasmModule => {
    essentia = new Essentia(wasmModule, false);
  });
}
```

```
async function fetchAudioAndDecode(url) {
  const audioBuffer = await essentia.getAudioBufferFromURL(url, audioCtx);
  const audioArray = essentia.audioBufferToMonoSignal(audioBuffer);
  return audioArray;
}
```

```
// analyse on button click
function analyse(audio) {
  // warning: avoid if using large audio file, analyse in chunks (framewise)
  let audioVector = essentia.arrayToVector(audio);
  const algoOutput = essentia.RMS(audioVector);

  audioVector.delete();

  return algoOutput;
}
```

# On Node.js

```
let esPkg = require("essentia.js");

// core essentia.js API
esPkg.Essentia
// essentia WASM backend
esPkg.EssentiaWASM
// add-on modules
esPkg.EssentiaModel
esPkg.EssentiaExtractor
esPkg.EssentiaPlot

// create an instance of Essentia core JS interface
const essentia = new esPkg.Essentia(esPkg.EssentiaWASM);
```

# ES6 import

```
import { EssentiaWASM } from './essentia-wasm.es.js';
import { Essentia } from './essentia.js-core.es.js';
import * as EssentiaModel from './essentia.js-model.es.js';
import * as EssentiaPlot from './essentia.js-plot.es.js';
import * as EssentiaExtractor from './essentia.js-extractor.es.js';

// create an instance of Essentia core JS interface
const essentia = new Essentia(EssentiaWASM);
```

Here, EssentiaWASM backend is loaded asynchronously



# Hands-on session

Go to: <https://glitch.com/@jmarcosfer/wac-21-essentia-js-tutorial>

We will cover:

- Non-realtime use
  - Main UI thread: <https://glitch.com/~essentiajs-core-non-rt>
  - Workers: <https://glitch.com/~essentiajs-core-non-rt-worker>
  - Node.js: <https://glitch.com/~essentiajs-core-node>
- Realtime use
  - ScriptProcessorNode: <https://glitch.com/~essentiajs-core-realtime-spn>
  - AudioWorklets: <https://glitch.com/~essentiajs-core-realtime-aw>

# Implementation Details

Some implementation details (mostly RT):

- Non-JS object clean-up:

```
let audioVector = essentia.arrayToVector(audio);  
const algoOutput = essentia.RMS(audioVector);  
  
audioVector.delete(); ←
```

- Code injection in AudioWorklets → ES6 imports unsupported for Worklets on Firefox
  - P. Adenot's [URLFromFiles](#) (thanks!)
  - used [here](#)
- Frame-size matching
  - using ChromeLabs [wasm-audio-helper.js](#)
- Getting analysis results from Worklet to UI thread with [SharedArrayBuffer](#)
  - P. Adenot's [ringbuf.js](#)

## Q/A Session

# Break time



# Machine Learning on the Web



ONNX

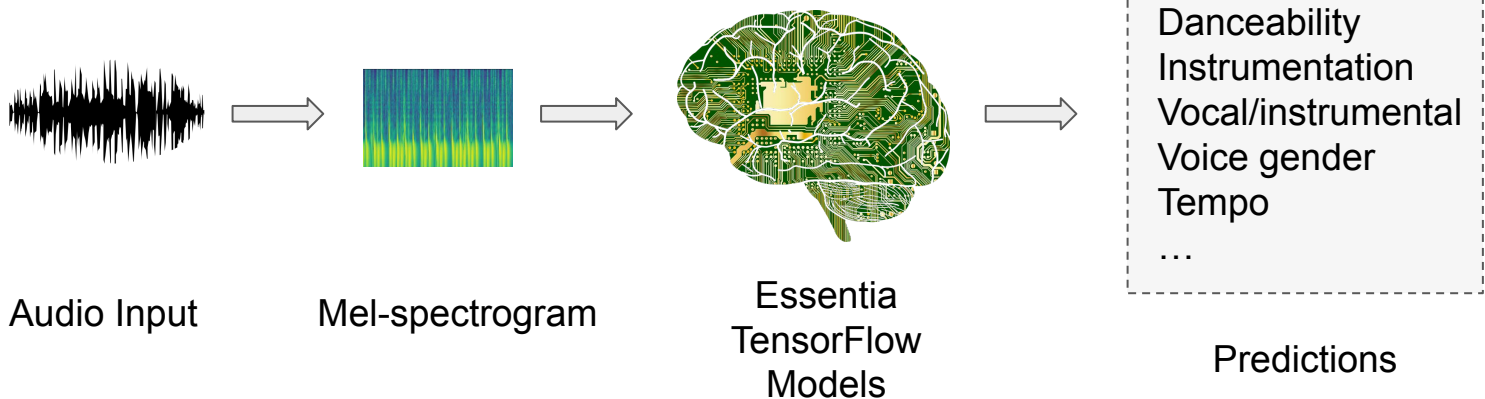


etc..

# Essentia TensorFlow models


A repository of publicly available deep learning models for music analysis tasks

[https://essentia.upf.edu/machine\\_learning.html](https://essentia.upf.edu/machine_learning.html)




Now, we ported these optimized models to Tensorflow.js format for its usage on JS.

# Auto-tagging and classification




*Queen - Bohemian Rhapsody*

# Auto-tagging and classification



Audio chunk  
(receptive field)

*Queen - Bohemian Rhapsody*




Predictions vector (activation values)




# Auto-tagging and classification


MSD tag activations



MTT tag activations



Classifier activations



*Queen - Bohemian Rhapsody*

# Auto-tagging / embeddings / tempo models

Model	RF (s)	Params.	Size (MB)	Purpose
MusiCNN	3	787K	3.1	AT/TL
VGG	3	605K	2.4	AT/TL
VGGish	1	62M	276	TL
TempoCNN	12	[27K-1.2M]	[0.1-4.7]	Tempo

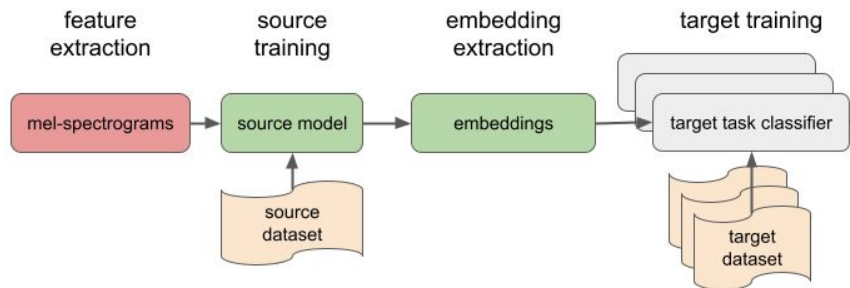
RF = receptive field

AT = auto-tagging

TL = transfer learning (embeddings)

# Transfer learning classifiers

## How we train the models



MusiCNN + MSD/MTT  
VGG + MSD/MTT  
VGGish + AudioSet

	Task	Classes
genre	dortmund	alternative, blues, electronic, folk-country, funksoulrnb, jazz, pop, raphiphop, rock
	gtzan	blues, classic, country, disco, hip hop, jazz, metal, pop, reggae, rock
	rosamerica	classic, dance, hip hop, jazz, pop, rhythm and blues, rock, speech
mood	acoustic	acoustic, non acoustic
	aggressive	aggressive, non aggressive
	electronic	electronic, non electronic
	happy	happy, non happy
	party	party, non party
	relaxed	relaxed, non relaxed
misc.	sad	sad, non sad
	danceability	danceable, non danceable
	voice/instrum.	voice, instrumental
	gender	male, female
	tonal/atonal	atonal, tonal
	urbansound8k	air conditioner, car horn, children playing, dog bark, drilling, engine idling, gun shot, jackhammer, siren, street music
	fs-loop-ds	bass, chords, fx, melody, percussion

# Essentia TensorFlow Models downloads

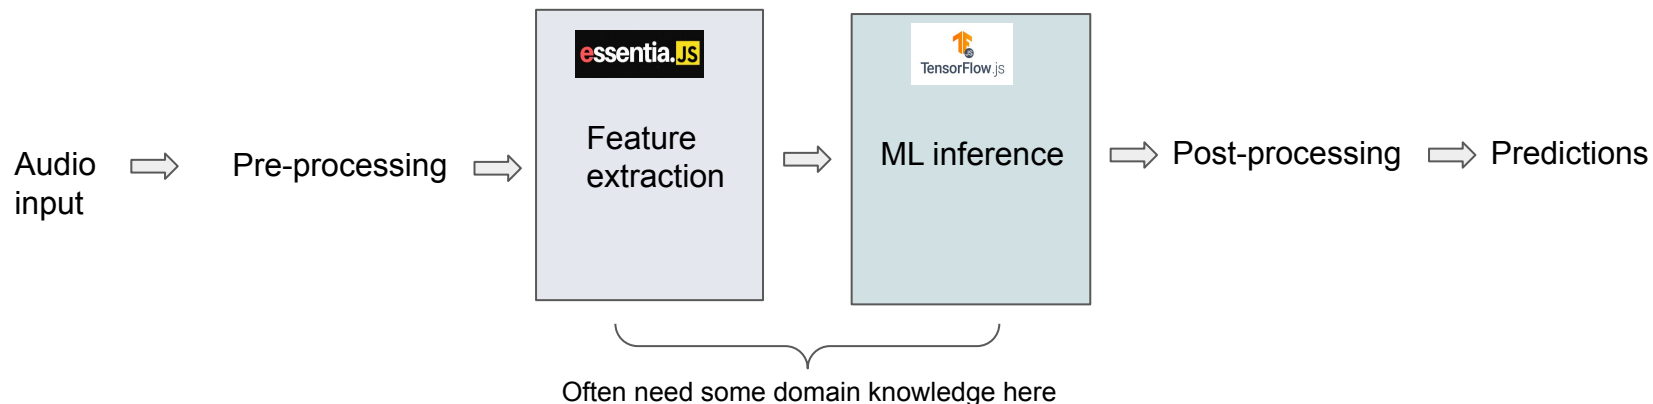
<https://essentia.upf.edu/models/>

[https://essentia.upf.edu/machine\\_learning.html](https://essentia.upf.edu/machine_learning.html)

Models available under [the CC BY-NC-ND 4.0 license](#)

# Essentia.js + TensorFlow.js

## Inference pipeline

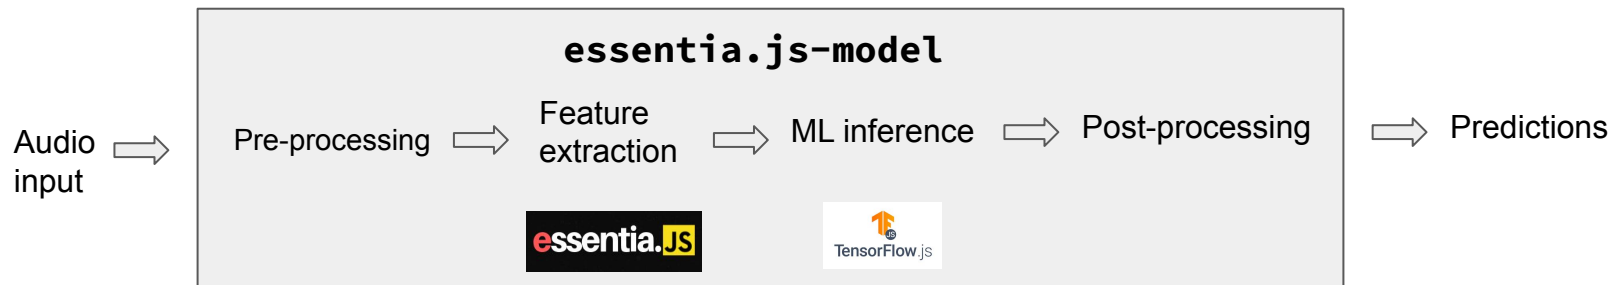


**But how to make this more accessible and easy to use for everyone?**

# essentia.js-model

- Add-on module for the essentia.js library.
- Simple JS API
- Support Web Workers for feature extraction and inference separately.

## Inference pipeline



[A. Correya, P. Alonso-Jiménez, J. Marcos-Fernández, X. Serra, and D. Bogdanov. Essentia TensorFlow models for audio and music processing on the web. In Web Audio Conference \(WAC 2021\), 2021.](#)

Let's checkout some models in action!

[Autotagging with MusiCNN](#)

[Mood Classification](#)

[illegible]

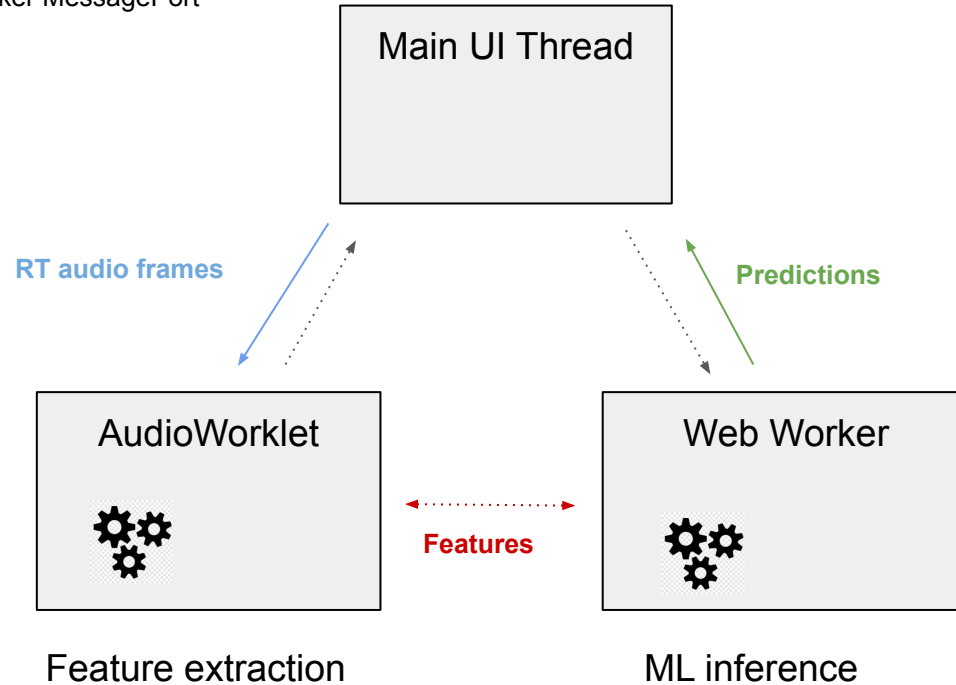


# Hands on

- Hands-on exercise using MusiCNN non-realtime: <https://glitch.com/~essentiajs-models-non-rt>
- Also can be used with Web Workers
- Realtime: <https://glitch.com/~essentiajs-models-rt>

# Real-time implementation

..... Worklet → Worker MessagePort



# Industrial application




## Audio Problems


Example: <https://glitch.com/~audio-problems-detection>

### Audio Problems


#### Incorrect margins




#### Phase/Stereo problems




#### Audio artifacts




#### Low frequency humming



#### Loudness problems



#### True-Peak detection



P. Alonso-Jiménez, L. Joglar-Ongay, X. Serra, and D. Bogdanov. Automatic Detection of Audio Problems for Quality Control in Digital Music Distribution. In AES Convention 146 (March 2019).

# Industrial application

## Custom extractor

- For better performance on JS
- Some algorithms are not yet working in Essentia.js



How to:

[https://github.com/MTG/essentia.js/blob/master/src/cpp/custom/essentia\\_custom\\_extractor.h](https://github.com/MTG/essentia.js/blob/master/src/cpp/custom/essentia_custom_extractor.h)  
[https://github.com/MTG/essentia.js/blob/master/src/cpp/custom/essentia\\_custom\\_extractor.cpp](https://github.com/MTG/essentia.js/blob/master/src/cpp/custom/essentia_custom_extractor.cpp)  
[https://github.com/MTG/essentia.js/blob/master/src/cpp/custom/bindings\\_extractor.cpp](https://github.com/MTG/essentia.js/blob/master/src/cpp/custom/bindings_extractor.cpp)

Then compile using the provided Makefile ensuring you have all requirements (or using the docker image)

# Industrial application

## Custom extractor

```
std::vector<float> SaturationDetectorExtractor::computeStarts(const val& audioData) {  
  
    std::vector<float> audioSignal = float32ArrayToVector(audioData);  
    std::vector<Real> frameFrameCutter;  
    std::vector<Real> startsSaturationDetector;  
  
    _FrameCutter->input("signal").set(audioSignal);  
    _FrameCutter->output("frame").set(frameFrameCutter);  
    _SaturationDetector->input("frame").set(frameFrameCutter);  
    _SaturationDetector->output("starts").set(startsSaturationDetector);  
  
    while (true) {  
        _FrameCutter->compute();  
        If (!frameFrameCutter.size()) {  
            break;  
        }  
        if (isSilent(frameFrameCutter)) continue;  
        _SaturationDetector->compute();  
    }  
    return startsSaturationDetector;  
};
```

# QA & Playground session

Build your own MIR web application :)

## Instructions

- Use any JS playground of your choice eg: Glitch, Runkit, Codepen, jsfiddle etc.
- You can use any algorithms, models or their combinations to make a fun prototype.
- Most welcome to remix our [examples on Glitch](#).

Optional,

- Tag your essentia.js prototypes on Twitter tagging @wac2021 with hashtag #essentiajs-wac2021

# Thanks

If you have any more questions, feedbacks etc,  
feel free to reach out to any of us

[@albincorreya](#)

[@di\\_bogdanov](#)

[@MaferGeorge](#)

[@luisjoqlar](#)



Universitat  
Pompeu Fabra  
Barcelona

**MTG**  
Music Technology  
Group